

# Parallel Programming Assignment 1

Ayush Raina, 22148

February 15, 2025

## Merge Sort

Merge Sort is a divide and conquer algorithm that divides the input array into two halves, recursively sorts the two halves, and then merges the sorted halves. The time complexity of Sequential Merge Sort is  $\mathcal{O}(n \log n)$ . We will implement some parallel versions of Merge Sort using CUDA C++ and Open MP.

## Benchmarking

We will use array size  $\sim 16.7$  million for initial observations. Later we will report observations with different array Sizes. Sequential Algorithm took  $\sim 4615$  ms = 4.615 sec to sort this array of size 16.7 million. We will use this as a reference to calculate speedup.

## Merge Sort Version 1

In this version, we will assume that division of array is done into single elements. We will start merging sorted sub lists of size 1, 2, 4, 8, .. and so on sequentially in parallel. Assume our array is denoted by  $\mathcal{A}$

- In Iteration 1, we will merge  $(\mathcal{A}[0], \mathcal{A}[1]), (\mathcal{A}[2], \mathcal{A}[3]), (\mathcal{A}[4], \mathcal{A}[5]), (\mathcal{A}[6], \mathcal{A}[7])$  and so on.
- In Iteration 2, we will merge  $(\mathcal{A}[0 - 1], \mathcal{A}[2 - 3]), (\mathcal{A}[4 - 5], \mathcal{A}[6 - 7])$  and so on.
- In Iteration 3, we will merge  $(\mathcal{A}[0 - 3], \mathcal{A}[4 - 7])$  and so on.

### Implementation 1

```
1  __global__ void ParallelMergeKernel(int* input, int* output, int size, int window_size) {
2
3      int tid = blockDim.x * blockIdx.x + threadIdx.x;
4      if(tid >= size) return;
5
6      // Calculate required indices for sublists
7      int start_one = tid * window_size * 2;
8      int end_one = min(start_one + window_size, size);
9
10     int start_two = end_one;
11     int end_two = min(start_two + window_size, size);
12
13     // Sequential Merging of two sublists
14     int i = start_one, j = start_two, k = start_one;
15
16     while(i < end_one and j < end_two) {
17         if(input[i] < input[j]) {
18             output[k++] = input[i++];
19         } else {
20             output[k++] = input[j++];
21         }
22     }
23
24     // Copy remaining elements in both cases
25     while(i < end_one) {
26         output[k++] = input[i++];
27     }
28
29     while(j < end_two) {
30         output[k++] = input[j++];
31     }
32 }
```

## Analysis of Version 1

In this version, we do not need  $N$  threads in total, we need:

- $$\text{threadsPerBlock} = \frac{N + 2 * \text{window\_size} - 1}{2 * \text{window\_size}}$$
- $$\text{blocksPerGrid} = \frac{N + \text{threadsPerBlock} - 1}{\text{threadsPerBlock}}$$

Here are the execution times for different blockSizes and windowSizes:

blockSize	Execution Time (ms)	Speedup
16	2046	2.2556
32	2048	2.2534
64	2049	2.2523
128	2047	2.2545
256	2050	2.2512

## Optimization 1: Binary Search Based Merging

In this version, we will use binary search to find the correct position of elements from one sublist in another sublist. Suppose  $pos = k$ , then first copy  $k - 1$  elements from first sublist to output array and then copy  $k$  elements from second sublist to output array and continue.

### Implementation 2

```
1  __global__ void ParallelMergeKernel(int* input, int* output, int size, int window_size) {
2      int tid = blockDim.x * blockIdx.x + threadIdx.x;
3      if(tid >= size) return;
4
5      // Calculate required indices for sublists
6      int start_one = (tid * 2) * window_size;
7      int end_one = min(start_one + window_size, size);
8
9      int start_two = end_one;
10     int end_two = min(start_two + window_size, size);
11
12     if(start_one >= size) return;
13
14     // Initialize pointers for output
15     int i = start_one;
16     int j = start_two;
17     int k = start_one;
18
19     // Binary search based merge
20     while(i < end_one && j < end_two) {
21
22         // Find position of input[j] in first subarray
23         if(i < end_one) {
24             int low = i, high = end_one, position = i;
25
26             while(low < high) {
27
28                 int mid = low + (high - low) / 2;
29                 if(input[mid] <= input[j]) {
30                     low = mid + 1;
31                     position = low;
32                 } else {
33                     high = mid;
34                 }
35             }
36
37             // Copy elements from first subarray up to the found position
38             while(i < position) {
39                 output[k++] = input[i++];
40             }
41
42             // Copy element from second subarray
43             output[k++] = input[j++];
44         }
45     }
46
47     // Copy remaining elements from first subarray
48     while(i < end_one) {
49
```

```

50     output[k++] = input[i++];
51 }
52
53 // Copy remaining elements from second subarray
54 while(j < end_two) {
55     output[k++] = input[j++];
56 }
57 }

```

## Analysis of Optimization 1

Here are the execution times for different blockSize and windowSizes:

blockSize	Execution Time (ms)	Speedup
16	1541	2.9948
32	1540	2.9967
64	1544	2.9889
128	1549	2.9793
256	1543	2.9909

## Merge Sort in Open MP

Here is the merge function which is mostly similar to sequential version:

### Implementation

```

1 void merge(vector<int>& arr, int left, int mid, int right) {
2     vector<int> temp(right - left + 1);
3     int i = left, j = mid + 1, k = 0;
4
5     while (i <= mid and j <= right) {
6         if (arr[i] <= arr[j])
7             temp[k++] = arr[i++];
8         else
9             temp[k++] = arr[j++];
10    }
11
12    while (i <= mid) {
13        temp[k++] = arr[i++];
14    }
15
16    while (j <= right) {
17        temp[k++] = arr[j++];
18    }
19
20    for (i = 0; i < k; i++) {
21        arr[left + i] = temp[i];
22    }
23
24 }

```

```

1 // Recursive merge sort function with OpenMP parallelization
2 // omp_set_num_threads = 24 is set in main function
3
4 void mergeSortParallel(vector<int>& arr, int left, int right, int depth = 0) {
5     if (left < right) {
6         int mid = left + (right - left) / 2;
7
8         // Parallelize only up to a certain depth to avoid overhead
9         if (depth < 4) {
10            #pragma omp parallel sections
11            {
12                #pragma omp section
13                mergeSortParallel(arr, left, mid, depth + 1);
14
15                #pragma omp section
16                mergeSortParallel(arr, mid + 1, right, depth + 1);
17            }
18        } else {
19            mergeSortParallel(arr, left, mid, depth + 1);
20            mergeSortParallel(arr, mid + 1, right, depth + 1);
21        }
22
23        merge(arr, left, mid, right);
24    }
25 }

```

## Analysis of Open MP Version

Here are the execution time for Open MP vs Sequential Merge Sort:

Array Size	Execution Time (ms)	Speedup
16.7 million	2481	1.8657

## Experimenting with Different Array Sizes with CUDA

Here are the execution times for different array sizes:

Array Size	Sequential (ms)	CUDA (ms)	Speedup (CUDA v1)	CUDA v2	Speedup (CUDA v2)
0.26 million	60	37	1.62	28	2.14
0.52 million	122	62	1.96	51	2.39
1.04 million	252	204	1.23	96	2.625
2.09 million	501	251	1.99	188	2.64
4.19 million	1014	510	1.98	384	1.8657
8.39 million	2085	1004	2.07	758	2.656
16.7 million	4290	2075	2.06	1562	2.746
33.5 million	8671	4061	2.13	3064	2.829
67.1 million	17754	8208	2.16	6165	2.879
134.2 million	36779	16040	2.292	12128	3.032
268.4 million	75431	32747	2.303	24584	3.068

## Experimenting with Different Array Sizes with Open MP

Here are the execution times for different array sizes:

Array Size	Execution Time (ms)	Speedup
0.26 million	31	1.93
0.52 million	65	1.87
1.04 million	137	1.83
2.09 million	278	1.80
4.19 million	572	1.77
8.39 million	1182	1.76
16.7 million	2437	1.76
33.5 million	5081	1.70
67.1 million	10373	1.71
134.2 million	21479	1.71
268.4 million	44450	1.69

## Plots

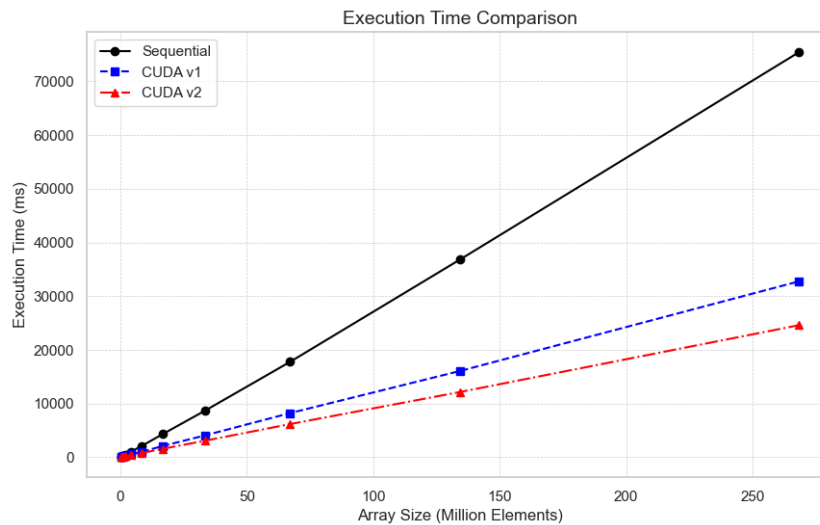


Figure 1: Execution Time vs Array Size for CUDA

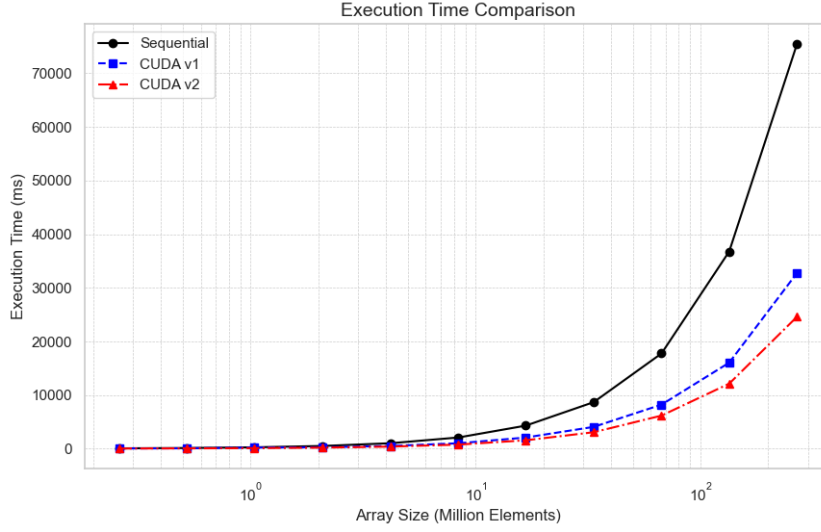


Figure 2: Execution Time vs Array Size for CUDA, x on log scale

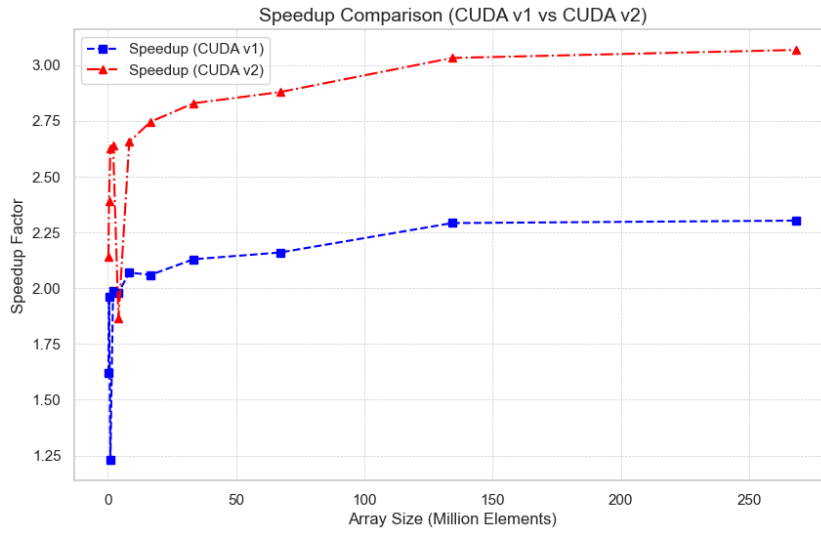


Figure 3: Speed up vs Array Size for CUDA

## 0.1 CUDA Kernel Profiling Results

Profiling was conducted using `nvprof` to analyze the execution performance of the CUDA kernel. The results are summarized as follows:

### 0.1.1 GPU Activities

The majority of GPU execution time was spent in the `ParallelMergeKernel` function:

- **ParallelMergeKernel:** 99.30% of total GPU time, with an average execution time of 140.56 ms per call.
- **Memory Transfers:**
  - [CUDA memcpy HtoD]: 0.36% of total GPU time, averaging 5.576 ms per call.
  - [CUDA memcpy DtoH]: 0.34% of total GPU time, averaging 5.289 ms per call.

### 0.1.2 CUDA API Calls

The CUDA API calls also reflect the kernel execution behavior:

- **cudaDeviceSynchronize()** accounted for 98.50% of total API execution time, with an average duration of 64.42 ms per call.
- **cudaMalloc()** took 0.78% of API execution time, averaging 6.13 ms per call.
- **cudaMemcpy()** accounted for 0.70%, with an average duration of 5.48 ms per call.




Figure 4: Speed up vs Array Size for CUDA, x on log scale



Figure 5: Execution Time vs Array Size for Open MP, x on log scale

### 0.1.3 Observations and Optimization Potential

- The `ParallelMergeKernel` function dominates execution time, suggesting that further optimization (such as improved memory access patterns or parallel workload balancing) may enhance performance.
- Memory transfers between host and device account for a small but non-negligible portion of execution time. Strategies like memory coalescing and asynchronous transfers could help reduce this overhead.
- Frequent calls to `cudaDeviceSynchronize()` indicate possible inefficiencies, as synchronization stalls execution until all GPU operations are complete.



Figure 6: Execution Time vs Array Size for Open MP, x and y on log scale




Figure 7: Speed up vs Array Size for Open MP



Figure 8: Comparison on Execution Times




Figure 9: Comparison on Speedup - x,y on log scale




Figure 10: Comparison on Speedup