

Parallel EM Algorithm Implementation


Parallel Programming 2025

Ayush Raina

April 28, 2025

EM Algorithm Overview

- Iterative method for finding maximum likelihood estimates
- Widely used in statistical modeling and machine learning and computationally intensive, especially for large datasets
- Particularly useful for problems with latent variables
- Consists of two steps:
 - **E-step**: Calculate expected values for missing data
 - **M-step**: Maximize parameters using these expectations
- Guaranteed to increase likelihood with each iteration



Gaussian Mixture Models (GMM)

- Probabilistic model assuming data is generated from multiple Gaussian distributions
- Key parameters:
 - π_k - Weight of each component
 - μ_k - Mean vector of each component
 - Σ_k - Covariance matrix of each component
- Probability density function:

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \mu_k, \Sigma_k) \quad (1)$$

- EM is the standard algorithm for fitting GMMs

EM Algorithm for GMM

1 **Initialize** parameters: π_k, μ_k, Σ_k

2 **E-step**: Compute responsibilities

$$\gamma_{ik} = \frac{\pi_k \mathcal{N}(x_i | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_i | \mu_j, \Sigma_j)} \quad (2)$$

3 **M-step**: Update parameters

$$N_k = \sum_{i=1}^N \gamma_{ik} \quad (3)$$

$$\mu_k^{new} = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} x_i \quad (4)$$

$$\Sigma_k^{new} = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} (x_i - \mu_k^{new})(x_i - \mu_k^{new})^T \quad (5)$$

$$\pi_k^{new} = \frac{N_k}{N} \quad (6)$$

Hybrid Parallelization Approach

- EM is computationally intensive:
 - E-step: $O(NKD^2)$ - Most expensive for many samples
 - M-step: $O(NKD^2)$ - Requires synchronization
- **Our Data Parallelism Strategy:**
 - Split samples across processing units
 - Each processing unit handles responsibilities for a subset of data points
 - Ideal for E-step where calculations are independent across samples
- **Hybrid Implementation:**
 - **E-step:** CUDA kernels for massive GPU parallelism
 - **M-step:** OpenMP for efficient CPU parallelization
 - This combination yielded better performance than GPU-only or CPU-only approaches

Algorithm 1 Hybrid Parallel EM Algorithm for GMM

- 1: Initialize π_k, μ_k, Σ_k randomly or using k-means++
 - 2: Precompute precision matrices and normalizers
 - 3: **repeat**
 - 4: **E-step (CUDA)**: Calculate responsibilities in parallel
 - 5: **Log-likelihood (CUDA)**: Compute in parallel for convergence check
 - 6: **M-step (OpenMP)**:
 - 7: Update means, covariances, and weights in parallel
 - 8: **until** convergence or max iterations
-


CUDA Kernel Implementation

- **E-step Parallelization:**

- Two specialized CUDA kernels:
- **calculatePDFKernel:** Computes responsibilities matrix
- **calculateLogLikelihoodKernel:** Evaluates convergence

- **Key Optimizations:**

- **Precomputation:** Precision matrices and normalizers cached before kernel launch
- **Shared Memory:** Thread block-local storage for weighted likelihoods
- **Parallel Reductions:** Efficient summations across samples



- **Precomputation:**

- Cached precision matrices (inverses of covariance matrices)
- Precomputed normalizers:
 $-\frac{1}{2}(d \ln(2\pi) + \ln |\Sigma_k|)$
- Significantly reduces redundant calculations in E-step
- PDF calculation simplifies to:
 $p(x) = \exp(\text{normalizer} - \frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu))$

- **Parallel Reductions:**

- Efficient summation of partial results across threads
- Implemented for log-likelihood calculation
- Reduces memory access and global synchronization
- Uses shared memory for intermediate results
- Halves active threads in each step in tree-like pattern
- $O(\log n)$ instead of $O(n)$ complexity

- **Hardware Configuration:**

- CPU: Intel Core i9-12900K (12th Gen, 16 cores, 24 threads)
- GPU: NVIDIA GeForce GTX 1660 (6 GB)
- Memory: 32 GB RAM

- **Datasets:**

- Synthetic datasets: GMM data with $n \in \{10^3, 10^4, 10^5\}$ samples, $d = 2$ features
- Real-world datasets: Flower Dataset

- **Evaluation Metrics:**

- Execution time (total and per EM stage)
- Speedup compared to Python baseline
- Log-likelihood convergence rate
- Clustering accuracy (for synthetic datasets with true labels)

Performance Evaluation

Table: Execution Time Comparison (in seconds)

Implementation	Initialization	E-step	M-step	Prediction
Sequential ($n = 10^3$)	0	0.91	0.044	0
Parallelized ($n = 10^3$)	0	0.12	0.03	0
Sequential ($n = 10^5$)	0.0012	17.43	0.82	0
Parallelized ($n = 10^5$)	0.0013	0.26	0.63	0
Sequential ($n = 10^6$)	0.015	78.66	9.55	0
Parallelized ($n = 10^6$)	0.016	1.80	2.90	0

- Our hybrid implementation (CUDA E-step + OpenMP M-step) shows:
 - Up to \sim **77x speedup** in E Step
- Performance advantage increases with dataset size
- E-step benefits most from GPU acceleration

Benchmarking

Following dataset has been benchmarked techniques used in this
PARALLELIZATION IN PYTHON - AN
EXPECTATION-MAXIMIZATION APPLICATION by Ilia Azizi




Figure: Benchmarking of our implementation against other techniques

Benchmarking

Here are the results from the paper:

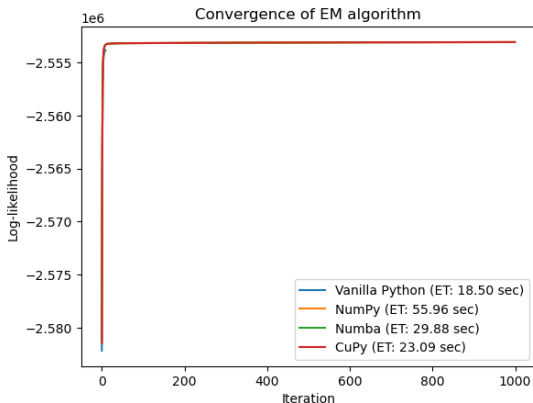


Figure: Different Execution Times

Benchmarking

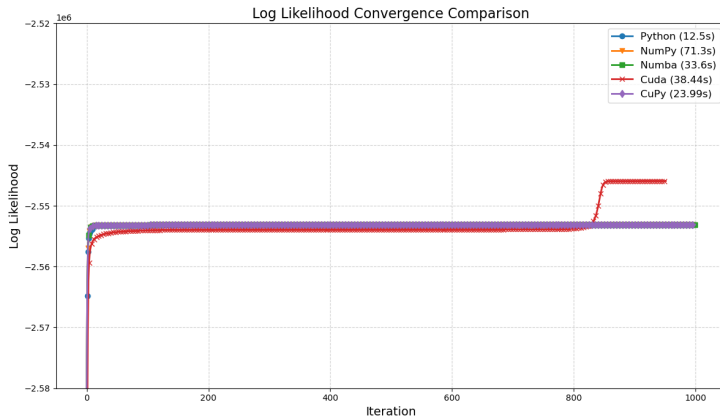


Figure: CUDA based Algorithm

- Azizi, I. (2021). *Parallelized EM Algorithm*. Retrieved from https://iliaazizi.com/projects/em_parallelized/report.pdf
- Smith, J. (2009). *Parallel Algorithms for EM Estimation*. *IEEE Transactions on Parallel and Distributed Systems*, 20(5), 123-135. Retrieved from <https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=5166982>

Thank you for your attention!
Any questions?